

Parallel and Concurrent Programming

Introduction and Foundation

Marwan Burelle

marwan.burelle@lse.epita.fr

<http://wiki-prog.infoprepa.epita.fr>

Parallel and Concurrent Programming Introduction and Foundation

Marwan Burelle

Introduction

Being Parallel

Foundations

Interacting with CPU Cache

CPU Cache

Mutual Exclusion

Definitions



1 Introduction

Parallelism in Computer Science

Nature of Parallelism

Global Lecture Overview

2 Being Parallel

Gain?

Models of Hardware Parallelism

Decomposition

3 Foundations

Tasks Systems

Program Determinism

Maximal Parallelism

Parallel and

Concurrent

Programming

Introduction and

Foundation

Marwan Burelle

Introduction

Being Parallel

Foundations

Interacting with

CPU Cache

Mutual Exclusion

Definitions

4 Interacting with CPU Cache

False Sharing
Memory Fence

Parallel and
Concurrent
Programming
Introduction and
Foundation
Marwan Burelle

5 Mutual Exclusion

Classic Problem: Shared Counter
Critical Section and Mutual Exclusion
Solutions with no locks

Introduction
Being Parallel
Foundations
Interacting with
CPU Cache
Mutual Exclusion
Definitions

6 Definitions

Parallel and
Concurrent
Programming
Introduction and
Foundation

Marwan Burelle

Introduction

Parallelism in Computer
Science

Nature of Parallelism

Global Lecture Overview

Being Parallel

Foundations

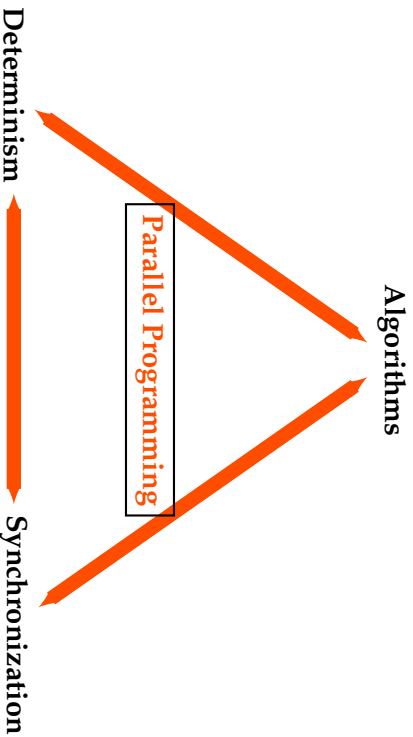
Interacting with
CPU Cache

Mutual Exclusion

Definitions

Introduction

- Next evolutions in processor tends more on more on growing of cores' number
- GPU and similar extensions follows the same path and introduce extra parallelism possibilities
- Network evolutions and widespread of internet fortify clustering techniques and grid computing.
- **Concurrency and parallelism are no recent concern, but are emphasized by actual directions of market.**



Parallel and
Concurrent
Programming
Introduction and
Foundation

Marwan Burelle

Introduction

Parallelism in Computer
Science

Nature of Parallelism

Global Lecture Overview

Being Parallel

Foundations

Interacting with
CPU Cache

Mutual Exclusion

Definitions

Parallelism in Computer Science

- **late 1950's:** first discussion about parallel computing.
- **1962:** *D825 by Burroughs Corporation* (four processors.)
- **1967:** Amdahl and Slotnick published a debat about feasibility of parallel computing and introduce *Amdahl's law* about limit of speed-up due to parallel computing.
- **1969:** *Honeywell's Multics* introduced first *Symmetric Multiprocessor* (SMP) system capable of running up to eight processors in parallel.
- **1976:** The first *Cray-1* is installed at *Los Alamos National Laboratory*. The major break-through of *Cray-1* is its *vector* instructions set capable of performing an operation on each element of a vector in parallel.
- **1983:** *CM-1 Connection Machine* by *Thinking Machine* offers 65536 1-bit processors working on a *SIMD* (Single Instruction, Multiple Data) model.

A Bit of History (2)

- **1991:** *Thinking Machine* introduced CM-5 using a MIMD architecture based on a fat tree network of SPARC RISC processors.
- **1990's:** modern micro-processors are often capable of being run in an SMP (Symmetric MultiProcessing) model. It began with processors such as *Intel's 486DX*, *Sun's UltraSPARC*, *DEC's Alpha* *IBM's POWER* ... Early SMP architectures was based on motherboard providing two or more sockets for processors.
- **2002:** *Intel* introduced the first processor with *Hyper-Threading* technology (running two threads on one physical processor) derived from DEC previous work.
- **2006:** First multi-core processors appears (several processors in one ship.)
- ...

How Can I Code In Parallel?

Processus Based

based No special support needed for parallelization but require complex inter-process communication. Not really useful.

System Level Threads

Using threads provides by system's libraries (like Unix *pthread*s). Lack of higher level features and portabilities issues.

Portable Threads Libs

reads Libs Solves portability issues of system threads.

Higher Level Libs

Can really increase efficiency and provide clever approach. May still need lower level interactions.

Language Level

Seriously ? Still very experimental.

Parallel air

Concurren

Programming

Introduction and

Foundation

Marwan Burelle

Introduction

Parallelism in Computer

Science

Nature of Parallelism

Global Lecture Overview

Being Parallel

Foundations

Interacting with

CPU Cache

Mutual Exclusion

Definitions

Parallel and
Concurrent
Programming
Introduction and
Foundation

Marwan Burelle

Introduction

Parallelism in Computer
Science

Nature of Parallelism

Global Lecture Overview

Being Parallel

Foundations

Interacting with
CPU Cache

Mutual Exclusion

Definitions

Nature of Parallelism

Parallel and
Concurrent
Programming
Introduction and
Foundation

Marwan Burelle

Introduction

Parallelism in Computer
Science

Nature of Parallelism

Global Lecture Overview

Being Parallel

Foundations

Interacting with

CPU Cache

Mutual Exclusion

Definitions

- **Vectorized Instructions:** usual extensions providing parallel execution of a single operation on all elements of a vector.
- **Non-Uniform Memory Access (NUMA):** architecture where each processor has its own memory with dedicated access (to limit false sharing for example.)

- In order to take advantage of hardware parallelism, the *operating system* **must** support SMP.
- *Operating Systems* can offer parallelism to application even when hardware is not, by using *multi-programming*.
- The system have two ways to provide parallelism: threads and processus.
- When only processus are available, no memory conflicts can arise, but processus have to use various communication protocol to exchange data and synchronized themselves.
- When memory is shared, programs have to manage memory accesses to avoid conflict, but data exchange between threads is far simpler.

- *Operating systems* supporting parallelism and threads, normally provide API for parallel programming.
 - Support for parallelism can be either primitive in the programming language or added by means of API and libraries.
 - The most common paradigm is the explicit use of threads backend by the Kernel.
 - Some languages try to offer implicit parallelism, or at least parallel blocks, but producing smart parallel code is tedious.
 - Some hardware parallelism features (vector instructions or multi-operations) may need special support from the language.
 - Modern frameworks for parallelism find a convenient way by abstracting system threads management and offering more simple threads manipulation (OpenMP, Intel's TBB ...)
- Parallel and Concurrent Programming Introduction and Foundation
- Marwan Burelle
- Introduction
Parallelism in Computer Science
Nature of Parallelism
Global Lecture Overview
Being Parallel
Foundations
Interacting with CPU Cache
Mutual Exclusion
Definitions

Parallel and
Concurrent
Programming
Introduction and
Foundation

Marwan Burelle

Introduction

Parallelism in Computer
Science

Nature of Parallelism

Global Lecture Overview

Global Lecture Overview

Being Parallel

Foundations

Interacting with
CPU Cache

Mutual Exclusion

Definitions

Parallel and
Concurrent
Programming
Introduction and
Foundation

Marwan Burelle

Introduction

Parallelism in Computer
Science

Nature of Parallelism

Global Lecture Overview

Being Parallel

Foundations

Interacting with

CPU Cache

Mutual Exclusion

Definitions

1 Introduction to parallelism

(this course)

2 Synchronization and Threads

How to enforce safe data sharing using various synchronization techniques, illustrated using Threads from C11/C++11.

3 Algorithms and Data Structures

How to adapt or write algorithms and data structures in a parallel world (shared queues, tasks scheduling, lock free structures ...)

4 TBB and other higher-level tools

Programming using Intel's TBB ...

Parallel and
Concurrent
Programming
Introduction and
Foundation

Marwan Burelle

Introduction

Being Parallel

Gain ?

Models of Hardware
Parallelism
Decomposition

Foundations

Interacting with
CPU Cache

Mutual Exclusion
Definitions

Being Parallel

Gain ?

Models of Hardware Parallelism

Foundations

Interacting with CPU Cache

Mutual Exclusion

Definitions

Parallel and
Concurrent
Programming
Introduction and
Foundation

Marwan Burelle

Amdahl's law

If P is a part of a computation that can be made parallel, then the maximum speed-up (with respect to the sequential version) of running this program on a N processors machine is:

$$\frac{1}{(1 - P) + \frac{P}{N}}$$

Introduction

Being Parallel

Gain ?

Models of Hardware
Parallelism

Decomposition

Foundations

Interacting with

CPU Cache

Mutual Exclusion

Definitions

Gustafson's law

Let P be the number of processor and α the sequential fraction of the parallel execution time, then we define the scaled-speedup, noted $S(P)$, as:

$$S(P) = P + \alpha \times (P - 1)$$

- Amdahl's law consider a fixed amount of work and focus on minimal time execution
- Gustafson's law consider a fixed execution time and describe the increased size of problem
- The main consequences of Gustafson's law is that, we can always increase size of the solved problem (for a fixed amount of time) by increasing the number of processor

Two Faces Of A Same Law

- It has been proved that both laws express the same result but in different form.
- Variation in results are often due to a misleading definition of Amdahl's P representing the parallelizable part of a sequential program and Gustafson's α representing the amount of time spent in the linear part during a parallel execution.

Parallel and Concurrent Programming Introduction and Foundation

Marwan Burelle

Introduction

Being Parallel

Gain ?

Models of Hardware

Parallelism

Decomposition

Foundations

Interacting with

CPU Cache

Mutual Exclusion

Definitions

Parallel and
Concurrent
Programming
Introduction and
Foundation

Marwan Burelle

Introduction

Being Parallel

Gain ?

Models of Hardware
Parallelism

Decomposition

Foundations

Interacting with
CPU Cache

Mutual Exclusion

Definitions

Models of Hardware Parallelism

	Single Instruction	Multiple Instruction
Single Data	SISD	MISD
Multiple Data	SIMD	MIMD

- **SISD**: usual non-parallel systems
- **SIMD**: performing the same operations on various data (like vector computing.)
- **MISD**: uncommon model where several operations are performed on the same data, usually implies that all operations must agreed on the result (fault tolerant code such as in space-shuttle controller.)
- **MIMD**: most common actual model.

Parallel and
Concurrent
Programming
Introduction and
Foundation

Marwan Burelle

Introduction

Being Parallel

Gain ?

Models of Hardware
Parallelism

Decomposition

Foundations

Interacting with
CPU Cache

Mutual Exclusion
Definitions

In Real Life?

- Actual processors provide **MIMD** parallelism in the form of Symmetric Multi-Processor (SMP)
- Modern processors also provides vector based instruction set extension that provides **SIMD** parallelism (like MMX or SSE instructions for x86)
- GPGPU are more or less used in a **SIMD** like fashion which is much more accurate for a very large number of core.

Parallel and Concurrent Programming Introduction and Foundation

Marwan Burelle

Introduction

Being Parallel

Gain?

Models of Hardware Parallelism

Decomposition

Foundations

Interacting with

CPU Cache

Mutual Exclusion

Definitions

Parallel and
Concurrent
Programming
Introduction and
Foundation

Marwan Burelle

Introduction

Being Parallel

Gain ?

Models of Hardware
Parallelism

Decomposition

Decomposition

Foundations

Interacting with
CPU Cache

Mutual Exclusion

Definitions

- To move from a linear *program* to a parallel *program*, we have to *break* our activities in order to things in parallel
- When decomposing you have to take into account various aspects:
 - Activities independance (how much synchronization we need)
 - Load Balancing (using all available threads)
 - Decomposition overhead

Choosing the right decomposition is the most critical choice you have to make when designing a parallel program

Decomposition Strategies



Task Driven

the problem is splitted in (almost) independent tasks ran in parallel;

Parallel and Concurrent Programming Introduction and Foundation

Introduction

Being Parallel

Gain ?

Models of Hardware Parallelism

Decomposition

Data Driven

all running task perform the same operations on a partition of the original data set;

Foundations

Interacting with

CPU Cache

Mutual Exclusion

Definitions

Data Flow Driven

the whole activities is decomposed in a chain of dependant tasks, a pipeline, where each task depends on the output of the previous one.

Task Driven Decomposition

- In order to perform task driven decomposition, you need to:
 - list activities in your program,
 - establish groups of dependent activities that will form standalone tasks,
 - identify interaction and possible data conflict between tasks
- Task driven decomposition is technically simple to implement and can be particularly efficient when activities are well segmented
- On the other hand, this strategy is highly constraint by the nature of your activities, their relationship and dependencies.
- Load balancing (maintaining thread activities at its maximum) is often hard to achieve due to the fixed nature of the decomposition.

- In order to perform data driven decomposition, you need to:
 - Defines the common task performed on each subset of data
 - Find a coherent data division strategy that respect load balancing, data conflict and other memory issue (such as cache false sharing)
 - You often have to prepare a recollection phase to compute final result (probably a fully linear computation,)
- Data driven decomposition scales pretty well, as data set size grows partitionning becomes more efficient against sequential or task based approach;
- Care must be taken when choosing data partitionning in order to obtain maximal performances
- Too much partitionning or too small data set will probably induce higher overhead.

- In order to perform data flow driven decomposition, you need to:
 - Split activities along the flow of execution in order to identify tasks
 - Model data exchange between tasks
 - Choose a data partitionning (the flow's grain) strategy
- With respect to Ford's concept of production line, execution time for a chunk of data correspond to the execution time of the longest task, the global time is thus this execution time multiply by the number of chunks plus two times the cost of the whole line;
- Needs carefull design and conception, data channels and efficient data partitionning, probably the more complex (but the more realistic) approach

The Choice of The Pragmatic Programmer

- Data driven approach yield pretty good result when performing single set of operations on huge set of data;
- Task driven approach are more suited for concurrency issues (performing various activities in parallel to mask waiting time.)
- Data flow driven decomposition can be used, together with another decomposition, as a skeleton for the global flow.
- Data flow driven decomposition requires complex infrastructure but provides a more adaptable solution for a complete chain of activities;
- Of course, in realistic situation, we'll often choose a mid-term decomposition mixing various approaches.

There are various design patterns (we'll see some later) dedicated to parallel programming, here are some examples:

- **Divide and Conquer:** data decomposition (divide) and recollection (conquer)
- **Pipeline:** data flow decomposition
- **Wave Front:** parallel topological traversal of a graph of tasks
- **Geometric Decomposition:** divide data-set in *rectangles*

Parallel and
Concurrent
Programming
Introduction and
Foundation

Marwan Burelle

Introduction

Being Parallel

Gain ?

Models of Hardware
Parallelism

Decomposition

Foundations

Interacting with

CPU Cache

Mutual Exclusion

Definitions

Parallel and
Concurrent
Programming
Introduction and
Foundation

Marwan Burelle

Introduction

Being Parallel

Foundations

Foundations

Tasks Systems
Program Determinism
Maximal Parallelism

Interacting with
CPU Cache

Mutual Exclusion

Definitions

Parallel and
Concurrent
Programming
Introduction and
Foundation

Marwan Burelle

Introduction

Being Parallel

Foundations

Tasks Systems

Program Determinism

Maximal Parallelism

Interacting with
CPU Cache

Mutual Exclusion

Definitions

Tasks Systems

- We will describe *parallel programs* by a notion of task.
- A task T is an instruction in our program. For the sake of clarity, we will limit our study to task of the form:

$T : \text{VAR} = \text{EXPR}$

where VAR is a memory location (can be seen as a variable) and EXPR are usual expressions with variables, constants and basic operators, but no function calls.

- A task T can be represented by two sets of memory locations (or variables): $\text{IN}(T)$ the set of memory locations used as input and $\text{OUT}(T)$ the set of memory locations affected by T .
- $\text{IN}(T)$ and $\text{OUT}(T)$ can, by them self, be seen as elementary task (as reading or writing values.) And thus our finest grain description of a program execution will be a sequence of $\text{IN}()$ and $\text{OUT}()$ tasks.

Example:

Let P_1 be a simple sequential program we present it here using task and memory locations sets:

T1 : $x = 1$	T1 : $\text{IN}(T1) = \emptyset$
T2 : $y = 5$	OUT(T1) = {x}
T3 : $z = x + y$	T2 : $\text{IN}(T2) = \emptyset$
T4 : $w = x - y $	OUT(T2) = {y}
T5 : $r = (z + w)/2$	T3 : $\text{IN}(T3) = \{x, y\}$
	OUT(T3) = {z}
	T4 : $\text{IN}(T4) = \{x, y\}$
	OUT(T4) = {w}
	T5 : $\text{IN}(T5) = \{z, w\}$
	OUT(T5) = {r}

- Given two sequential programs (a list of tasks) a parallel execution is a list of tasks resulting of the composition of the two programs.
- Since, we do not control the scheduler, the only constraint on an execution is the preservation of the order between tasks of the same program.
- Scheduling does not *undestand* our notion of task, it rather works at assembly instructions level, and thus, we can assume that a task T can be interleaved with another task between the realisation of the subtask $IN(T)$ and the realisation of the subtask $OUT(T)$.
- As for tasks, the only preserved order is that $IN(T)$ always appears before $OUT(T)$.
- Finally, an execution can be modeled by an ordered sequence of input and output sets of memory locations.

Parallel and
Concurrent
Programming
Introduction and
Foundation
Marwan Burelle

Introduction
Being Parallel

Foundations

Tasks Systems

Program Determinism
Maximal Parallelism

Interacting with
CPU Cache

Mutual Exclusion

Definitions

Execution (*example*)

Example:

Given the two programs P_1 and P_2 :

$T11 : x = 1$

$T21 : y = 1$

$T12 : y = x + 1$

$T22 : x = y - 1$

The following sequences are valid parallel execution of $P_1 // P_2$:

$E1 = \text{IN}(T11); \text{OUT}(T11); \text{IN}(T12); \text{OUT}(T12); \text{IN}(T21); \text{OUT}(T21); \text{OUT}(T22)$
 $E2 = \text{IN}(T21); \text{OUT}(T21); \text{IN}(T22); \text{OUT}(T22); \text{IN}(T11); \text{OUT}(T11); \text{IN}(T12); \text{OUT}(T12)$
 $E3 = \text{IN}(T11); \text{IN}(T21); \text{OUT}(T11); \text{OUT}(T21); \text{IN}(T12); \text{IN}(T22); \text{OUT}(T22); \text{OUT}(T12)$

At the end of each executions we can observe each value in both memory locations x and y :

$E1 \quad x = 0 \text{ and } y = 1$

$E2 \quad x = 1 \text{ and } y = 2$

$E3 \quad x = 0 \text{ and } y = 2$

- Parallel and Concurrent Programming Introduction and Foundation
- Marwan Burelle
- Introduction
- Being Parallel
- Foundations
- Tasks Systems
 - Program Determinism
 - Maximal Parallelism
 - Interacting with CPU Cache
 - Mutual Exclusion
- Definitions

Parallel and
Concurrent
Programming
Introduction and
Foundation

Marwan Burelle

Introduction

Being Parallel

Foundations

Tasks Systems

Program Determinism

Maximal Parallelism

Interacting with
CPU Cache

Mutual Exclusion

Definitions

Program Determinism

Tasks' Dependencies

- In order to completely describe parallel programs and parallel executions of programs, we introduce a notion of dependencies between tasks.
- Let E be a set of tasks and $(<)$ a *well founded dependency order* on E .
- A pair of tasks T_1 and T_2 verify $T_1 < T_2$ if the sub-task $\text{OUT}(T_1)$ **must** occurs before the sub-task $\text{IN}(T_2)$.
- A **Task System** $(E, <)$ is the definition of a set, E , of tasks and a dependency order $(<)$ on E . It describes a combination of several sequential programs into a parallel program (or a fully sequential program if $(<)$ is total.) Tasks of a same sequential program have a *natural* ordering, but we can also define ordering between tasks of different programs, or between programs.

Task Language

Task Language


Let $E = \{T_1, \dots, T_n\}$ be a set of task, $A = \{\text{IN}(T_1), \dots, \text{OUT}(T_n)\}$ a vocabulary based on sub-task of E and ($<$) an ordering relation on E .

The language associated with a task system $S = (E, <)$, noted $L(S)$, is the set of words ω on the vocabulary A such that for every T_i in E there is exactly one occurrence of $\text{IN}(T_i)$ and one occurrence of $\text{OUT}(T_i)$ and the former appearing before the latter. If $T_i < T_j$ then $\text{OUT}(T_i)$ must appear before $\text{IN}(T_j)$.

- We can define the product of system S_1 and S_2 by $S_1 \times S_2$ such that $L(S_1 \times S_2) = L(S_1).L(S_2)$ (._._ is the concatenation of language.)
- We can also define parallel combination of task system: $S_1 // S_2 = (E_1 \cup E_2, <_1 \cup <_2)$ (where $E_1 \cap E_2 = \emptyset$.)

Precedence Graph (example)

```
{  
  T1;  
  ///  
  {T2;T3;T4};  
  {T5;T6;T7};  
  ///  
  T8;  
}
```



If we define $S1 = \{T1\}$, $S2 = \{T2\ T3\ T4\}$, $S3 = \{T5\ T6\ T7\}$ and $S4 = \{T8\}$. Then the resulting system (described by the graph above) is:

$$S = S1 \times (S2 / S3) \times S4$$

Deterministic System

A deterministic task system $S = (E, <)$ is such that for every pair of words ω and ω' of $L(S)$ and for every memory locations X , sequences values affected to X are the same for ω and ω' .

A deterministic system, is a tasks system where every possible executions are not distinguishable by only observing the evolution of values in memory locations (observational equivalence, a kind of bisimulation.)

Parallel and
Concurrent
Programming
Introduction and
Foundation

Marwan Burelle

Introduction

Being Parallel

Foundations

Tasks Systems

Program Determinism

Maximal Parallelism

Interacting with

CPU Cache

Mutual Exclusion

Definitions

- The previous definition may seem *too restrictive* to be useful.
- In fact, one can exclude *local* memory locations (*i.e.* memory locations not shared with other programs) of the observational property.
- In short, the deterministic behavior can be limited to a restricted set of meaningful memory locations, excluding temporary locations used for inner computations.
- The real issue here is the *provability* of the deterministic behavior: one can not possibly test every execution path of a given system.
- We need a finite property independant of the scheduling (*i.e.* a property relying only on the system.)

- Non-Interference (*NI*) is a general property used in many context (especially language level security.)
- Two tasks are non-interfering, if and only if the values taken by memory locations does not depend on the order of execution of the two tasks.

Non Interference

Let $S = (E, <)$ be a tasks system, T_1 and T_2 be two task of E , then T_1 and T_2 are non-interfering if and only if, they verify one of the two following properties:

- $T_1 < T_2$ or $T_2 < T_1$ (the system force a particular order.)
- $\text{IN}(T_1) \cap \text{OUT}(T_2) = \text{IN}(T_2) \cap \text{OUT}(T_1) = \text{OUT}(T_1) \cap \text{OUT}(T_2) = \emptyset$

- The NI definitions is a based on the contraposition of the Bernstein's conditions (defining when two tasks are dependent.)
- Obviously, two non-interfering tasks do not introduce non-deterministic behavior in a system (they are already ordered or the order of their execution is not relevant.)

Theorem

Let $S = (E, <)$ be a tasks system, S is a deterministic system if every pair of tasks in E are non-interfering.

- We now extend our use of observational equivalence to compare systems.
- The idea is that we can not distinguish two systems that have the same behavior (affect the same sequence of values in a particular set of memory locations.)

Equivalent Systems

Let $S_1 = (E_1, <_1)$ and $S_2 = (E_2, <_2)$ be two tasks systems. S_1 and S_2 are equivalent if and only if:

- $E_1 = E_2$
- S_1 and S_2 are deterministic
- For every words $\omega_1 \in L(S_1)$ and $\omega_2 \in L(S_2)$, for every (meaningful) memory location X , ω_1 and ω_2 affect the same sequence of values to X .

Parallel and
Concurrent
Programming
Introduction and
Foundation

Marwan Burelle

Introduction

Being Parallel

Foundations

Tasks Systems

Program Determinism

Maximal Parallelism

Maximal Parallelism

Interacting with
CPU Cache

Mutual Exclusion

Definitions

Maximal Parallelism

- Now that we can define and verify determinism of tasks systems, we need to be able to assure a kind of maximal parallelism.
- Maximal parallelism describes the minimal sequentiality and ordering needed to stay deterministic.
- A system with maximal parallelism can't be *more* parallel without introduction of non-deterministic behavior (and thus inconsistency.)
- Being able to build (or transform systems into) maximally parallel systems, guarantees usage of a *parallel-friendly* computer at its maximum capacity for our given solution.

Maximal Parallelism

Maximal Parallelism

A tasks system with maximal parallelism, is a tasks where one can not remove dependency between two tasks T_1 and T_2 without introducing interference between T_1 and T_2 .

Theorem

For every deterministic system $S = (E, <)$ there exists an equivalent system with maximal parallelism $S_{max} = (E, <_{max})$ with ($<_{max}$) defined as:

$$T_1 <_{max} T_2 \text{ if } \left\{ \begin{array}{l} T_1 <_{\emptyset} T_2 \\ \wedge OUT(T_1) \neq \emptyset \wedge OUT(T_2) \neq \emptyset \\ \wedge \left(\begin{array}{l} IN(T_1) \cap OUT(T_2) \neq \emptyset \\ \vee IN(T_2) \cap OUT(T_1) \neq \emptyset \\ \vee OUT(T_1) \cap OUT(T_2) \neq \emptyset \end{array} \right) \end{array} \right.$$

- Given a graph representing a system, one can reason about parallelism and performances.
- Given an (hypothetical) unbound material parallelism, the complexity of a parallel system is the length of the longest path in the graph from initial tasks (tasks with no predecessors) to final tasks (tasks with no successor.)
- Classical analysis of dependency graph can be use to spot critical tasks (tasks that can't be late without slowing the whole process) or find good planned excutions for non-parallel hardware.
- Tasks systems and maximal parallelism can be used to prove modelization of parallel implementations of sequential programs.
- Maximal parallelism can be also used to effectively measure real gain of a parallel implementation.

Parallel and
Concurrent
Programming
Introduction and
Foundation

Marwan Burelle

Introduction

Being Parallel

Foundations

Interacting with
CPU Cache

False Sharing

Memory Fence

Mutual Exclusion

Definitions

Interacting with CPU Cache

Cache: Hidden Parallelism Nightmare

- Modern CPUs rely on memory cache to prevent memory access bottleneck;
- In SMP architecture, cache are mandatory: there's only one memory bus ! (NUMA architectures try to solve this)
- Access to shared data induce memory locking and cache updates (thus waiting time for your core.)
- Even when data are not explicitly shared, cache management can become your worst enemy !

This part is Intel/x86 oriented, technical details may not correspond to other processors.

- Each cache line have a special state: **Invalid(I)**, **Shared(S)**, **Exclusive(E)**, **Modified(M)**
- Cache line in state S are shared among core and only used for reading.
- Cache line in state E or M are only owned by one core.
- When a core tries to write to some memory location in a cache line, it will forced any other core to *loose* the cache line (putting it to state I for example.)
- Cache mechanism worked as a read/write lock that track modifications and consistency between values in cache line and value in memory.
- The pipeline (and somehow the compiler) try to *anticipate* write needs so cache line are directly acquired in E state (*fetch for write*)

Parallel and
Concurrent
Programming
Introduction and
Foundation
Marwan Burelle

Introduction
Being Parallel

Foundations

Interacting with
CPU Cache

False Sharing
Memory Fence

Mutual Exclusion

Definitions

False Sharing

Parallel and
Concurrent
Programming
Introduction and
Foundation

Marwan Burelle

Introduction

Being Parallel

Foundations


Interacting with
CPU Cache

False Sharing

Memory Fence

Mutual Exclusion

Definitions



There's no standard ways to control cache interaction, even using ASM code. Issues described previously may or may not happen in various contexts. Rather than providing a seminal solutions, we must rely on guidelines to prevent cache false sharing.

- Avoid as much as possible shared data;
- Prefer *threads' local storages* (local variable, locally allocated buffers...);
- When returning set of values, allocate a container per working thread;
- Copy shared data before using it;
- When possible use a thread oriented allocator (modern allocator will work with separated pools of memory per unit rather than one big pool);

Even Better

```
int main()
{
    int          *res[2];
    pthread_t     t[2];
    // Provide no containers
    pthread_create(t, NULL, run, NULL);
    pthread_create(t+1, NULL, run, NULL);
    // let threads allocate the result
    // and collect it with join.
    pthread_join(t[0], res);
    pthread_join(t[1], res+1);
}
```

Parallel and
Concurrent
Programming
Introduction and
Foundation
Marwan Burelle

Introduction

Being Parallel

Foundations


Interacting with
CPU Cache

False Sharing
Memory Fence

Mutual Exclusion

Definitions

Some stats ...



Input (n)	Bad Practice	Good Practice	Difference
1024	0.0446	0.0028	0.0418
2048	0.1878	0.0107	0.177078
4096	0.7140	0.0423	0.671683
8192	2.7296	0.1688	2.560812
16384	7.4191	0.6817	6.73742
65536	122.0350	10.7911	111.2439
131072	682.3750	43.1707	639.2043

The bad code take more then 90% longer, due to false sharing.

Both code do heavy computation based on n , bad version share an array between threads for reading and writing, while the other copy the input value and allocate its own container. Time results are in seconds, measured using `tbb::tick_count`. The main program runs 4 threads on 2 dual core processor (no hyperthreading) under FreeBSD.

Parallel and
Concurrent

Programming
Introduction and
Foundation

Marwan Burelle

Introduction

Being Parallel

Foundations

Interacting with

CPU Cache

False Sharing

Memory Fence

Mutual Exclusion

Definitions

Don't Trust The Evidence

- Modern processor are able to somehow modify execution order.
- On multi-processor platform this means that apparent ordering may not be respected at execution level (in fact, your compiler is doing the same)

Form Intel's TBB Documentation

Another mistake is to assume that conditionally executed code cannot happen before the condition is tested. However, the compiler or hardware may speculatively hoist the conditional code above the condition.

Similarly, it is a mistake to assume that a processor cannot read the target of a pointer before reading the pointer. A modern processor does not read individual values from main memory. It reads cache lines . . .

Parallel and Concurrent Programming Introduction and Foundation

Marwan Burelle

Introduction

Being Parallel

Foundations

Interacting with

CPU Cache

False Sharing

Memory Fence

Mutual Exclusion

Definitions

Trap

You fool ...

```
bool Ready;
std::string Message;

// Thread 1 action
void Send( const std::string& src ) {
    Message=src; // C++ hidden memory
    Ready = true;
}

// Thread 2 action
bool Receive( std::string& dst ) {
    bool result = Ready;
    if( result ) dst=Message; // C++ hidden memory
    return result;
}
```


Parallel and
Concurrent
Programming
Introduction and
Foundation

Marwan Burelle

Introduction

Being Parallel

Foundations

Interacting with
CPU Cache

Mutual Exclusion

Classic Problem: Shared
Counter
Critical Section and Mutual
Exclusion
Solutions with no locks

Definitions

Mutual Exclusion

Parallel and
Concurrent
Programming
Introduction and
Foundation

Marwan Burelle

Introduction

Being Parallel

Foundations

Interacting with
CPU Cache

Mutual Exclusion

Classic Problem: Shared
Counter

Critical Section and Mutual
Exclusion

Solutions with no locks

Definitions

Classic Problem: Shared Counter

- Sharing a counter between two threads (or processes) is good seminal example to understand the complexity of synchronisation.
- The problem is quite simple: we have two threads monitoring external events, when an event occurs they increase a global counter.

- Increasing a counter X is a simple task of the form:

T1 : $X = X + 1$

With associated sets:

$IN(T1) = \{X\}$

$OUT(T1) = \{X\}$

- The two thread execute the same task T1 (along with their monitoring activity.) And thus, they are interfering.

A shared counter without locking.

Threads	Average Errors	Ratio
2	3878947.2	46.24%
4	11575434.6	68.99%
8	28256366	84.21%
16	57525557.4	85.72%
32	119176589.8	88.79%
64	244048804	90.92%
128	478563124.4	89.14%
256	965023132.4	89.87%

Programs run on a 4-cores (8 hyper-threads) CPU. Each threads (system threads using C++11 API) wait for everyone and then perform a 2^{22} steps loop where it adds one to the counter at each step.

Parallel and
Concurrent
Programming
Introduction and
Foundation

Marwan Burelle

Introduction

Being Parallel

Foundations

Interacting with
CPU Cache

Mutual Exclusion

Classic Problem: Shared
Counter

Critical Section and Mutual
Exclusion

Solutions with no locks

Definitions

Pseudo-Code for Counter Sharing

Example:

```
global int X=0;

guardian(int id):
    for (;;)
        wait_event(id); // wait for an event
        X = X + 1;      // T1

main:
    { // // parallel execution
        guardian(0);
        guardian(1);
    }
```

Foundations Interacting with

Marwan Burelle

Being Parallel

Interacting with CPU Cache

Classic Problem: Shared Counter

Solutions with no locks

Definitions

- In the previous example, while we can easily see the interference issue, no task ordering can solve it.
- We need an other refinement to enforce consistency of our program.
- The critical task (T1) is called a **critical section**.

Critical Section

A section of code is said to be a critical section if execution of this section can not be interrupted by other process manipulating the same shared data without loss of consistency or determinism.

The overlapping portion of each process, where the shared variables are being accessed.

Example:

```
guardian(int id):  
    // Restant Section  
    for (;);  
    wait_event(id);  
    // Entering Section (empty here)  
    X = X + 1; // Critical Section  
    // Leaving Section (empty here)
```

Restant Section

: section outside of the critical part

Critical Section (CS)

: section manipulating shared data

Entering Section

: code used to enter CS

Leaving Section

: code used to leave CS

- | |
|--|
| Parallel and Concurrent Programming Foundation |
| Marwan Burelle |
| Introduction |
| Being Parallel |
| Foundations |
| Interacting with CPU Cache |
| Mutual Exclusion |
| Classic Problem: Shared Counter |
| Critical Section and Mutual Exclusion |
| Solutions with no locks |
| Definitions |

- **Deadlock:** two process try to enter in CS at the same time and block each other (errors in *entering section*.)
- **Race condition:** two processes make an assumption that will be invalidated when the execution of the other process will finished (no mutual exclusion.)
- **Starvation:** a process is waiting for CS indefinitely.
- **Priority Inversion:** a complex double blocking situation between process of different priority. In short, a high priority process is consuming execution time waiting for CS, blocking a low priority process already in CS (spin waiting.)

Parallel and Concurrent Programming Introduction and Foundation

Marwan Burelle

Introduction

Being Parallel

Foundations

Interacting with CPU Cache

Mutual Exclusion

Classic Problem: Shared Counter
Critical Section and Mutual Exclusion
Exclusion

Solutions with no locks

Definitions

Solutions with no locks

Example:

```
global int X=0;
global int turn=0;

guardian(int id):
    int other = (id+1)%2;
    for (;;)
        wait_event(id);
        while(turn!=id);
        X = X + 1;
        turn=other;
```

- This solution enforce mutual exclusion: turn cannot have two different values at the same time.
- This solution enforce bounded waiting: you can see the other thread passing only one time while waiting.
- This solution **does not respect progression**:
 - You will wait for entering **CS** that the other thread is passed (even if it arrived after you.)
 - If the other thread see no event, it will not go through the **CS** and won't let you take your turn !

Example:

```
global int X=0;
global int ASK[2] = {0;0};

guardian(int id):
  int other = (id+1)%2;
  for (;;)
    wait_event(id);
    ASK[id] = 1;
    while(ASK[other]);
    X = X + 1;
    ASK[id] = 0;
```

Parallel and
Concurrent
Programming
Introduction and
Foundation

Marwan Burelle

Introduction

Being Parallel

Foundations

Interacting with
CPU Cache

Mutual Exclusion

Classic Problem: Shared
Counter

Critical Section and Mutual
Exclusion

Solutions with no locks

Definitions

- This solution enforce mutual exclusion: turn cannot have two different values at the same time.
- This solution respects progression
- This solution **present a dead lock**:
 - When asking for **CS**, each thread will set their flag and then waits if necessary
 - Both thread can set their flag simultaneously
 - Thus, both thread will wait each other with escape possibility

Example:

```
global int X=0;
global int ASK[2] = {0;0};

guardian(int id):
int other = (id+1)%2;
for (;;)
    wait_event(id);
    while(ASK[other]);
    ASK[id] = 1;
    X = X + 1;
    ASK[id] = 0;
```

Parallel and
Concurrent
Programming
Introduction and
Foundation

Marwan Burelle

Introduction

Being Parallel

Foundations

Interacting with
CPU Cache

Mutual Exclusion

Classic Problem: Shared
Counter

Critical Section and Mutual
Exclusion

Solutions with no locks

Definitions

- This tiny modification removed the dead lock of previous solution
- But, this solution **present a race condition, mutual exclusion is violated:**
 - When entering **CS**, a thread will first wait and then set its flag
 - Both thread can enter the waiting loop **before** the other one has set its flag and then just pass
 - Both thread can thus enter the **CS**: **lost game** !

The Peterson's Algorithm

Example:

```
global int X=0;
global int turn=0;
global int ASK[2] = {0;0};

guardian(int id):
    int other = (id+1)%2;
    for (;;)
        wait_event(id);
        ASK[id] = 1;
        turn=other;
        while(turn!=id && ASK[other]);
        X = X + 1;
        ASK[id] = 0;
```

The Peterson's Algorithm

- The previous algorithm satisfies *mutual exclusion*, *progress* and *bounded waiting*.
- The solution is limited to two process but can be generalized to any number of processes.
- This solution is *hardware/system independent*.
- The main issue is *spin wait*: a process waiting for CS is consuming time resources, opening risks of *priority inversion*.

Parallel and
Concurrent
Programming
Introduction and
Foundation

Marwan Burelle

Introduction

Being Parallel

Foundations

Interacting with
CPU Cache


Mutual Exclusion

Classic Problem: Shared
Counter

Critical Section and Mutual
Exclusion

Solutions with no locks

Definitions



Parallel and
Concurrent
Programming
Introduction and
Foundation

Marwan Burelle

Introduction

Being Parallel

Foundations

Interacting with
CPU Cache

Mutual Exclusion

Definitions

Definitions

Dependency Ordering Relation

a dependency ordering relation is a partial order which verifies:

- anti-symmetry ($T_1 < T_2$ and $T_2 < T_1$ can not be both true)
- anti-reflexive (we can't have $T < T$)
- transitive (if $T_1 < T_2$ and $T_2 < T_3$ then $T_1 < T_3$).

◀ back

Parallel and
Concurrent
Programming
Introduction and
Foundation

Marwan Burelle

Introduction
Being Parallel

Foundations

Interacting with
CPU Cache

Mutual Exclusion

Definitions

Task Language

Task Language

Let $E = \{T_1, \dots, T_n\}$ be a set of task, $A = \{\text{IN}(T_1), \dots, \text{OUT}(T_n)\}$ a vocabulary based on sub-task of E and ($<$) an ordering relation on E .

The language associated with a task system $S = (E, <)$, noted $L(S)$, is the set of words ω on the vocabulary A such that for every T_i in E there is exactly one occurrence of $\text{IN}(T_i)$ and one occurrence of $\text{OUT}(T_i)$ and the former appearing before the latter. If $T_i < T_j$ then $\text{OUT}(T_i)$ must appear before $\text{IN}(T_j)$.

◀ back

► [back](#)

Parallel and
Concurrent
Programming
Introduction and
Foundation

Marwan Burelle

Introduction
Being Parallel

Foundations

Interacting with
CPU Cache

Mutual Exclusion

Definitions

Transitive Closure

The transitive closure of a relation ($<$) is the relation $<_{\mathcal{C}}$ defined by:

$$x <_{\mathcal{C}} y \text{ if and only if } \left\{ \begin{array}{l} x < y \\ \exists z \text{ such that } x <_{\mathcal{C}} z \text{ and } z <_{\mathcal{C}} y \end{array} \right.$$

This relation is the biggest relation that can be obtained from ($<$) by only adding sub-relation by transitivity.

◀ back

Parallel and
Concurrent
Programming
Introduction and
Foundation

Marwan Burelle

Introduction

Being Parallel

Foundations

Interacting with

CPU Cache

Mutual Exclusion

Definitions

Equivalent Relation

Let $<$ be a well founded partial order, any relation $<_{eq}$ is said to be equivalent to $<$ if and only if, $<_{eq}$ has the same transitive closure as $<$.

Kernel ($<_{min}$)

The kernel $<_{min}$ of a relation $<$ is the smallest relation equivalent to $<$, that is if we suppress any pair $x <_{min} y$ to the relation it is no longer equivalent.

► back

Parallel and
Concurrent
Programming
Introduction and
Foundation

Marwan Burelle

Introduction

Being Parallel

Foundations

Interacting with

CPU Cache

Mutual Exclusion

Definitions

Deterministic System

A deterministic task system $S = (E, <)$ is such that for every pair of words ω and ω' of $L(S)$ and for every memory locations X , sequences values affected to X are the same for ω and ω' .

A deterministic system, is a tasks system where every possible executions are not distinguishable by only observing the evolution of values in memory locations (observational equivalence, a kind of bisimulation.)

◀ back

Non Interference

Let $S = (E, <)$ be a tasks system, T_1 and T_2 be two task of E , then T_1 and T_2 are non-interfering if and only if, they verify one of the two following properties:

- $T_1 < T_2$ or $T_2 < T_1$ (the system force a order.)
- $\text{IN}(T_1) \cap \text{OUT}(T_2) = \text{IN}(T_2) \cap \text{OUT}(T_1) = \text{OUT}(T_1) \cap \text{OUT}(T_2) = \emptyset$

► [back](#)

Parallel and Concurrent Programming Introduction and Foundation

Marwan Burelle

Introduction

Being Parallel

Foundations

Interacting with

CPU Cache

Mutual Exclusion

Definitions

Equivalent Systems

Equivalent Systems

Let $S_1 = (E_1, <_1)$ and $S_2 = (E_2, <_2)$ be two tasks systems. S_1 and S_2 are equivalent if and only if:

- $E_1 = E_2$
- S_1 and S_2 are deterministic
- For every words $\omega_1 \in L(S_1)$ and $\omega_2 \in L(S_2)$, for every (meaningful) memory location X , ω_1 and ω_2 affect the same sequence of values to X .

► [back](#)

Parallel and Concurrent Programming Introduction and Foundation

Marwan Burelle

Introduction

Being Parallel

Foundations

Interacting with

CPU Cache

Mutual Exclusion

Definitions

Maximal Parallelism

Maximal Parallelism

A tasks system with maximal parallelism, is a tasks where one can not remove dependency between two tasks T_1 and T_2 without introducing interference between T_1 and T_2 .

► back

Theorem

For every deterministic system $S = (E, <)$ there exists an equivalent system with maximal parallelism $S_{\max} = (E, <_{\max})$ with $(<_{\max})$ defined as:

$$T_1 <_{\mathcal{E}} T_2$$

$$T_1 <_{max} T_2 \text{ if } \left\{ \begin{array}{l} T_1 <_{\mathcal{E}} T_2 \\ \wedge OUT(T_1) \neq \emptyset \wedge OUT(T_2) \neq \emptyset \\ \wedge \left(\begin{array}{l} IN(T_1) \cap OUT(T_2) \neq \emptyset \\ \vee IN(T_2) \cap OUT(T_1) \neq \emptyset \\ \vee OUT(T_1) \cap OUT(T_2) \neq \emptyset \end{array} \right) \end{array} \right.$$

► back